

Classification of psychiatry clinical notes by diagnosis: A deep learning and machine learning approach

Sergio Rubio-Martín^{Corresp., 1}, María Teresa García-Ordás¹, Antonio Serrano-García², Clara Margarita Franch-Pato², Arturo Crespo-Álvaro¹, José Alberto Benítez-Andrade¹

¹ ALBA Research Group, Department of Electric, Systems and Automatics Engineering, Universidad de León, León, Spain

² Servicio de Psiquiatría, Complejo Asistencial Universitario de León (CAULE), León, Spain

Corresponding Author: Sergio Rubio-Martín
Email address: srubm@unileon.es

The classification of clinical notes into specific diagnostic categories is critical in healthcare, especially for mental health conditions like Anxiety and Adjustment Disorder. In this study, we compare the performance of various Artificial Intelligence models, including both traditional Machine Learning approaches (Random Forest, Support Vector Machine, K-nearest neighbors, Decision Tree, and eXtreme Gradient Boost) and Deep Learning models (DistilBERT and SciBERT), to classify clinical notes into these two diagnoses. Additionally, we implemented three oversampling strategies: No Oversampling, Random Oversampling, and Synthetic Minority Over-sampling Technique (SMOTE), to assess their impact on model performance. Hyperparameter tuning was also applied to optimize model accuracy. Our results indicate that oversampling techniques had minimal impact on model performance overall. The only exception was SMOTE, which showed a positive effect specifically with BERT-based models. However, hyperparameter optimization significantly improved accuracy across the models, enhancing their ability to generalize and perform on the dataset. The Decision Tree and eXtreme Gradient Boost models achieved the highest accuracy among machine learning approaches, both reaching 96%, while the DistilBERT and SciBERT models also attained 96% accuracy in the deep learning category. These findings underscore the importance of hyperparameter tuning in maximizing model performance. This study contributes to the ongoing research on AI-assisted diagnostic tools in mental health by providing insights into the efficacy of different model architectures and data balancing methods.

Classification of Psychiatry Clinical Notes by Diagnosis: A Deep Learning and Machine Learning Approach

Sergio Rubio-Martín¹, María Teresa García-Ordás¹, Antonio Serrano-García², Clara Margarita Franch-Pato², Arturo Crespo-Álvaro¹, and José Alberto Benítez-Andrade¹

¹ ALBA Research Group, Department of Electric, Systems and Automatics Engineering, Universidad de León, Campus of Vegazana s/n, León, 24071, León, Spain

² Servicio de Psiquiatría, Complejo Asistencial Universitario de León (CAULE), 24008, León, Spain

Corresponding author:

Sergio Rubio-Martín¹

Email address: srubm@unileon.es

ABSTRACT

The classification of clinical notes into specific diagnostic categories is critical in healthcare, especially for mental health conditions like Anxiety and Adjustment Disorder. In this study, we compare the performance of various Artificial Intelligence models, including both traditional Machine Learning approaches (Random Forest, Support Vector Machine, K-nearest neighbors, Decision Tree, and eXtreme Gradient Boost) and Deep Learning models (DistilBERT and SciBERT), to classify clinical notes into these two diagnoses. Additionally, we implemented three oversampling strategies: No Oversampling, Random Oversampling, and Synthetic Minority Over-sampling Technique (SMOTE), to assess their impact on model performance. Hyperparameter tuning was also applied to optimize model accuracy. Our results indicate that oversampling techniques had minimal impact on model performance overall. The only exception was SMOTE, which showed a positive effect specifically with BERT-based models. However, hyperparameter optimization significantly improved accuracy across the models, enhancing their ability to generalize and perform on the dataset. The Decision Tree and eXtreme Gradient Boost models achieved the highest accuracy among machine learning approaches, both reaching 96%, while the DistilBERT and SciBERT models also attained 96% accuracy in the deep learning category. These findings underscore the importance of hyperparameter tuning in maximizing model performance. This study contributes to the ongoing research on AI-assisted diagnostic tools in mental health by providing insights into the efficacy of different model architectures and data balancing methods.

1 INTRODUCTION

The field of medicine has advanced rapidly in recent decades due to technological innovations that have transformed both the diagnostic and treatment phases. In the mental health sector, particularly in psychiatry, there has been a paradigm shift, with a growing focus on understanding the brain and the underlying mechanisms that regulate behavior, emotions, and responses to external or internal changes. Despite these advances, psychiatric diagnosis still faces significant challenges due to the subjective and complex nature of the symptoms presented by patients, as well as the frequent overlap between different mental disorders. Moreover, despite the progress made in the field of medicine, the number of patients suffering from mental disorders has not decreased but has instead increased since 2019, when 970 million people were living with a mental disorder (World Health Organization, 2022a). The World Health Organization (WHO) is concerned not only about these numbers but also about the increase in mental disorders diagnosed during the COVID-19 pandemic, with cases of anxiety disorder rising by an estimated 26% (World Health Organization, 2022b).

In the field of mental healthcare, two major disciplines coexist: psychology and psychiatry. Although

46 they share the common goal of improving mental well-being, they differ significantly in their training,
47 methods, and approaches to treatment. Psychology is the scientific study of mental processes and behavior,
48 including both internal mental activities, such as thoughts and emotions, and externally observable
49 behaviors (Henriques and Michalski, 2020). Psychological practice primarily involves therapeutic
50 methods based on dialogue and behavioral interventions, such as cognitive-behavioral therapy, humanistic
51 therapy, or psychodynamic approaches. These treatments focus on modifying dysfunctional behaviors,
52 emotions, and thoughts, typically following a non-medical model.

53 Psychiatry, by contrast, is a branch of medicine concerned with the diagnosis, treatment, and prevention
54 of mental, emotional, and behavioral disorders. Psychiatrists, as medical doctors, are trained to assess
55 both psychological symptoms and their biological underpinnings. They can prescribe pharmacological
56 treatments and often manage complex cases involving severe mental illnesses, such as schizophrenia,
57 bipolar disorder, major depression, or severe anxiety disorders (Kendler et al., 2011).

58 While psychology and psychiatry approach mental health from different perspectives—psychology
59 focusing more on psychological and social aspects, psychiatry integrating biological, psychological, and
60 pharmacological considerations—the two disciplines are complementary and increasingly collaborate in
61 interdisciplinary mental health teams to provide holistic patient care.

62 While psychological conditions often involve significant distress, psychiatric disorders may pose more
63 serious risks to patients' lives, including an increased risk of suicide. Moreover, it has been extensively
64 documented that individuals suffering from severe mental disorders frequently experience reduced life
65 expectancy. For example, people diagnosed with schizophrenia have an estimated life expectancy that
66 is 10 to 20 years shorter than that of the general population (Nimavat et al., 2023). In addition, many
67 individuals with mental illnesses face a substantial treatment gap, with only 29% of those with psychosis
68 and 33% of those with depression receiving formal mental health care (World Health Organization, 2021;
69 Moitra et al., 2022). These challenges highlight the pressing need for innovative approaches to support
70 mental healthcare systems and improve access and quality of care.

71 One of the biggest problems for people suffering from a mental illness when seeking help from a public
72 psychologist or psychiatrist is the long waiting time to get an appointment. Timely access to professionals
73 would help patients receive a diagnosis and appropriate treatment for their condition. However, due to the
74 lack of human and economic resources, as well as the time required to get an appointment, we propose a
75 solution aimed at reducing the workload in the classification of clinical notes, also known as electronic
76 health records (EHR).

77 It has been shown that artificial intelligence (AI) models have helped in various medical fields. For
78 example, in oncology, AI has become a valuable tool for predicting cancer (Liu et al., 2020; Alanazi, 2023;
79 Briganti and Le Moine, 2020). Additionally, AI continues to be useful in predicting cancer recurrence
80 (Zhang et al., 2023). Some AI models based on images have been used to detect different types of cancer,
81 such as skin, breast, and lung cancer (Midasala et al., 2024; Kaka et al., 2022; Quanyang et al., 2024). In
82 other medical fields, AI has contributed significantly to improving outcomes, such as in the detection of
83 diabetes in patients (Wu, 2024).

84 However, AI models are not limited to using a single type of input, such as images; they can also
85 process text as a source of information. Natural language processing (NLP) techniques help extract
86 meaningful information from different types of texts. Among the goals of NLP is predicting whether
87 a person suffers from a particular illness. This approach has been applied, for example, in using AI to
88 predict whether a patient has autism spectrum disorder (ASD), achieving nearly 90% accuracy (Rubio-
89 Martín et al., 2024). Another study related to AI and psychiatry involved the classification of texts
90 about eating disorders (ED) into four categories—texts written by someone with ED, texts that promote
91 ED, informative texts, and scientific texts—achieving nearly 87% accuracy in one of the categories
92 (Benítez-Andrade et al., 2022).

93 Delving specifically into the convergence between psychiatry and AI, several studies have attempted
94 to assist in the diagnosis or classification of complex mental disorders, such as schizophrenia, depression,
95 or anxiety disorders, using AI (Kodipalli and Devi, 2023; Cortes-Briones et al., 2022; ALSAGRI and
96 YKHLEF, 2020; Nemeshure et al., 2021). As shown, applying NLP techniques can help extract relevant
97 information from unstructured data, such as EHRs. The use of EHRs as input for AI has led to the
98 development of models capable of predicting depression crises in patients (Msosa et al., 2023).

99 In recent years, these efforts have been further expanded in multiple directions. For instance, the
100 prediction of anxiety symptoms in social anxiety disorder has been achieved using multimodal data

101 collected during virtual reality sessions (Park et al., 2025). In another line of work, deep learning models
102 have been developed that outperform clinicians in identifying violence risk from emergency department
103 notes (Dobbins et al., 2024). Transformer-based models have also been employed to detect personal
104 and family history of suicidal ideation in EHRs, yielding F1-scores above 0.90 (Adekkanattu et al.,
105 2024). Furthermore, suicide risk has been phenotyped using multi-label classification strategies based on
106 psychiatric clinical notes (Li et al., 2024).

107 One of the most challenging scenarios in AI-driven classification involves EHRs, where patients are
108 diagnosed with various mental disorders that share overlapping symptoms. The differentiation between
109 anxiety disorders (ICD-10 F41) and adjustment disorders (ICD-10 F43) is key in the clinical diagnosis
110 and appropriate treatment of patients. Both disorders can present anxious symptoms, but these play a
111 different role in each case. In anxiety disorders (F41), anxious symptoms are central and form part of the
112 core clinical picture. Examples of these disorders include generalized anxiety disorder and social anxiety
113 disorder. Anxiety in these cases does not require a specific external triggering event; that is, the person
114 may experience excessive and ongoing worries about various aspects of life without a clear precipitating
115 factor (World Health Organization, 2019).

116 On the other hand, adjustment disorders (F43) are characterized by the presence of an identifiable life
117 event or stressor that triggers the symptoms, which may include anxiety, depression, or behavioral changes.
118 These symptoms are a disproportionate psychological response to a stressful situation, such as the loss
119 of a loved one, divorce, or work-related difficulties, and they are time-limited. Unlike anxiety disorders,
120 symptoms in adjustment disorders tend to resolve when the individual adjusts to the new situation or the
121 stressful event is resolved.

122 While anxiety disorders present anxious symptoms as a central element and do not rely on a clear
123 external trigger, adjustment disorders always have an identifiable stressful event that precipitates the
124 symptoms. This differentiation is fundamental to guide both diagnosis and therapeutic decisions. The
125 importance of distinguishing between these two types of disorders is crucial to avoid misdiagnosis, as
126 clinical interventions for each may differ significantly. A misdiagnosis or confusion between the two
127 could lead to inappropriate treatments, negatively affecting the patient's prognosis (Casey and Bailey,
128 2011).

129 For classification purposes, we grouped all ICD-10 codes under the F41 category (Other anxiety
130 disorders) into a single "anxiety disorder" class. This includes panic disorder or episodic paroxysmal
131 anxiety (F41.0), generalized anxiety disorder (F41.1), mixed anxiety disorders (F41.3), other specified
132 anxiety disorders (F41.8), and unspecified anxiety disorder (F41.9). Although our approach focuses
133 on analyzing and classifying existing clinical notes rather than intervening during the initial diagnostic
134 process, structuring and interpreting this information has substantial value. Enhanced documentation
135 quality, retrospective clinical audits, improved training datasets for future models, and support for research
136 activities are some of the ways in which structured clinical information can meaningfully contribute to the
137 mental healthcare system without altering the core diagnostic workflows.

138 Due to the challenges involved in classifying these two mental disorders, this research demonstrates
139 how AI can achieve highly accurate classification of EHRs, specifically aiming to identify patients
140 diagnosed with adjustment disorder or anxiety disorder. Additionally, this manuscript presents several
141 substantial advancements. The key contributions of this research include:

- 142 • **Machine learning models:** We trained several machine learning models in pursuit of the best
143 results. To optimize the performance of each model, a hyperparameter tuning process was carried
144 out. The implementation of this tuning process helped to improve the initial results.
- 145 • **BERT-based models:** We explored BERT models, testing two separate pretrained versions, each
146 with distinct training datasets and features that influenced their effectiveness in our tasks.
- 147 • **Data balancing process:** Although the dataset is sufficiently large to evaluate the metrics of
148 each model, we applied two data balancing techniques, known as Random Oversampling and
149 Synthetic Minority Oversampling Technique (SMOTE). These techniques were used to assess
150 whether increasing the number of samples in the dataset would allow the models to leverage
151 additional characteristics that could improve the classification task.
- 152 • **Real medical dataset:** For this research, clinical notes were provided by the 'Complejo Asisten-
153 cial Universitario de León' (CAULE). This dataset contains electronic health records of patients

154 diagnosed with adjustment and anxiety disorders. The dataset is entirely self-created, giving it
155 unique value and relevance. From its initial design and data collection to its cleaning, preprocessing,
156 and transformation, every step was meticulously handled to align with the goals of this research.
157 By controlling the entire data treatment process, we gained a deep understanding of the dataset's
158 structure, limitations, and potential insights. This level of control allows for highly tailored analyses
159 and more reliable results. Due to the challenges and restrictions in obtaining clinical notes or other
160 patient information, this dataset holds significant scientific value.

161 The paper is organized as follows: Section 2, 'Material and Methods', provides a detailed descrip-
162 tion of the methodology applied, including the collection and preprocessing of the dataset. Section 3,
163 'Experiments and Results', outlines the experiments conducted and presents the findings, along with a
164 comparison of the various models used. Lastly, Section 4, 'Discussion and Conclusions', brings together
165 the discussion and conclusion to create a unified narrative.

166 2 MATERIAL AND METHODS

167 This section provides a detailed explanation of the methodology implemented throughout the research.
168 Firstly, section 2.1 describes the process followed to obtain the dataset and how it was transformed from
169 unstructured to structured data. Next, section 2.2 presents the models implemented for this research.
170 Additionally, section 2.3 outlines the hardware and software specifications of the computer used for the
171 research.

172 2.1 Dataset collection and classification

173 All research involving patient information requires time and the ability to overcome several challenges that
174 arise throughout the process. To begin with, patients' EHRs contain highly sensitive information, which
175 must be protected under strict privacy regulations, as mandated by the European Union's General Data
176 Protection Regulation (GDPR) (European Parliament, 2016). Since patient identification is not required
177 for this research, the clinical notes were anonymized to allow the use of EHRs as a valuable information
178 source, not only in the medical field but also in the field of artificial intelligence (Rao et al., 2023).

179 An ethics committee was convened and granted us permission to use Spanish EHRs as a dataset for
180 research purposes, ensuring that no patient could be identified. This approval was issued by the Research
181 Ethics Committee for Medicinal Products of the Health Areas of León and Bierzo under the identifier
182 2303. The EHRs consist of clinical notes from the psychiatry unit of CAULE, written entirely in plain
183 text without any structured data. The dataset comprises 12,921 clinical notes, collected between January
184 11, 2017, and December 31, 2022. All clinical notes were collected from the Psychiatry Emergency
185 Service of the hospital. Each note documents an urgent psychiatric assessment performed during an
186 emergency department visit. These notes are not part of scheduled outpatient consultations or longitudinal
187 inpatient records, but rather correspond to acute episodes requiring immediate attention. Depending on
188 the evaluation, the patient is either discharged (often with referral for outpatient follow-up) or admitted to
189 inpatient care. Therefore, each note is self-contained and not part of a progressive sequence of visits.

190 This research was supported by professional psychiatrists who assisted in creating structures to
191 organize the information found in the EHRs. Additionally, these experts provided several guidelines for
192 processing the data. The first step in dataset preprocessing was to remove samples or records where the
193 clinical note was either empty or not properly completed.

194 To avoid including clinical notes that lacked sufficient or valuable information due to their short
195 length, the experts decided not to consider clinical notes shorter than 600 characters. This threshold
196 was not arbitrary but carefully determined, as it was found that many samples under 600 characters
197 lacked the necessary information to begin structuring the data. Moreover, it was calculated that applying
198 this threshold retained almost 95% of the dataset while ensuring that no clinical notes with insufficient
199 information were included, as shown in Figure 1.

200 Continuing with the preprocessing phase, the first data extracted from the EHRs were the patient's
201 age and gender. To achieve this, regular expressions were used. A preview of the dataset revealed various
202 patterns that allowed for the extraction of most patients' ages. All phrases structured like '20 years
203 old man' and '30 years old woman' among other possibilities, were captured using a complex regular
204 expression.




Figure 1. Distribution of the mean clinical note length (in characters) per patient. The green area shows that 95% of the patients have an average note length above the 600-character threshold applied. CIPA_AN are the patients ID.

Moreover, the regular expressions were designed to account for common human writing errors, such as missing or extra spaces between words, as well as misspelled words, ensuring the correct extraction of the patient's age. The task of extracting the patient's gender was partially accomplished using the same regular expression, as 'Man' and 'Woman' directly refer to male and female, respectively.

However, in some cases, extracting gender is more challenging, such as in clinical notes where the term 'Patient' is used instead of gender-specific terms like 'Man' or 'Woman'. In these instances, past participle verb forms in Spanish were used to infer the patient's gender. Additionally, when these verbs were absent, marital status indicators like 'single' or 'married', which have gender-specific forms in Spanish, were leveraged to help determine the patient's gender.

The new dataset now consists of several columns. The first column contains the original clinical note. The second column contains the patient's gender, represented as 'V' for male and 'M' for female. The third column records the patient's age. Since the psychiatric clinical notes are plain texts written by professionals summarizing the interview with the patient, the EHRs try to follow the Subjective-Objective-Assessment-Progress (SOAP) standard. However, in this dataset, the information for each section is not clearly delineated, and the majority of notes are composed as unstructured narratives rather than strictly segmented reports.

As a result, identifying the actual diagnosis from these notes is not straightforward. Diagnostic terms such as "anxiety" or "adjustment disorder" may appear in different parts of the note — for instance, in the personal or family history, in symptom descriptions, or as part of comorbidities — without necessarily representing the primary diagnosis. Additionally, anxiety is frequently recorded as a symptom within broader diagnostic categories, adding semantic ambiguity. For these reasons, we did not remove diagnostic terms from the clinical notes during preprocessing. This choice was deliberate, as our aim was to evaluate whether the model could correctly infer the diagnosis based on context, even in the presence of potentially misleading or overlapping terms.

The initial goal of this approach was to extract diagnoses from each clinical note. To achieve this, a Large Language Model (LLM), specifically ChatGPT 4.0, was utilized, as it has proven to be a powerful tool for information extraction in various research studies (Wang et al., 2023). For this research, the ChatGPT API, accessed through Microsoft Azure services, was employed to process 1,000 clinical notes. Prompt engineering techniques, including the use of different roles in API requests (García-Barragán

²³⁴ et al., 2024), were applied to enhance the model's performance.




Figure 2. Representation of the prompt structure.

²³⁵ One such technique was Few-Shot learning, which involves assigning the model a specific role,
²³⁶ explaining the task objective, breaking the task into multiple steps, and providing correct examples of
²³⁷ how the task should be performed. This approach ensures that the model understands how to execute the
²³⁸ task effectively. Several scientific publications emphasize the value of this Few-Shot prompt engineering
²³⁹ technique when working with ChatGPT. In this case, the ChatGPT 4.0 model, which can handle up to
²⁴⁰ 32,000 tokens of conversational context, was used. Since the clinical notes are written in Spanish, the
²⁴¹ prompt was constructed in Spanish; however, for ease of understanding in this paper, the prompt will be
²⁴² presented in English. The API request format is shown in Figure 2. The prompt structure is explained
²⁴³ below:

- ²⁴⁴ • **ROLE:** The role assigned to the model. This instruction helps the model adopt an appropriate
²⁴⁵ perspective, focusing on knowledge relevant to the designated role. In this case, the role given was:
²⁴⁶ 'You are an assistant and a linguist specialized in identifying entities within text. You are a leading
²⁴⁷ expert in psychiatry, and I need your help with a very important task in medicine'.
- ²⁴⁸ • **TASK and INSTRUCTIONS:** The objective of the task is explained to the model, outlining how it
²⁴⁹ should proceed and detailing how special situations should be handled. Furthermore, the process
²⁵⁰ is broken down into a list of instructions that can be easily followed by the model, as the main
²⁵¹ problem is split into smaller, manageable tasks.
- ²⁵² • **CLINICAL NOTE 1:** Corresponds to the first clinical note provided as plain text.
- ²⁵³ • **CORRECT ANSWER GIVEN:** A sample of a correct answer is provided to the model for the first
²⁵⁴ clinical note. This example helps the model understand how to proceed. In this case, it was specified
²⁵⁵ that the model should label the diagnosis as 'DX' during entity extraction, using '@ @' to indicate
²⁵⁶ the start of the extraction and '# #' to indicate the end of the diagnosis extraction. One example of a
²⁵⁷ correct answer given would be 'DX @@ Ansiedad reactiva, Sindrome ansioso-depresivo # #'.
- ²⁵⁸ • **CLINICAL NOTE 2:** The next clinical note provided to the model to continue the task.

²⁵⁹ The final results provided by the LLM were reviewed by experts. After completing the entire
²⁶⁰ preprocessing process, we focused on those clinical notes where patients were diagnosed with Adjustment
²⁶¹ Disorder or any form of Anxiety Disorder. For this line of research, which centers on these two mental




Figure 3. Demographic data from the patients found in the clinical notes.

| DX | Man(%) | Unknown(%) | Woman(%) | Mean (Years) | Std (Years) |
|---------------|--------|------------|----------|--------------|-------------|
| Adjustment D. | 32.9 | 6.2 | 60.9 | 44.4 | 18.3 |
| Anxiety D. | 36.6 | 2.4 | 61.0 | 42.7 | 16.4 |

Table 1. Percentage distribution by gender and complete age statistics by diagnosis. (Std = Standard Deviation)

disorders, a total of 228 clinical notes were considered: 82 corresponding to Adjustment Disorder and 146 to Anxiety Disorder.

As shown in the left part of Figure 3, of these 228 EHRs, it was found that 61% and 34.2% correspond to clinical notes where the patient is a woman and a man, respectively. Only 4.8% of the notes correspond to cases where the patient's sex is not identified. Additionally, the figure presents the age distribution of patients across the clinical notes, categorized in 5-year intervals. This histogram reveals that the majority of patients fall within the 30 to 50-year age range, with a notable peak around the age of 40. Specifically, the highest number of clinical notes corresponds to patients aged between 35 and 45 years. The distribution also shows that there are fewer clinical notes for patients below 20 years of age and above 70 years, indicating that the majority of the patient population receiving treatment for Adjustment Disorder and Anxiety Disorder tends to be middle-aged. This age trend is consistent with research that shows a high prevalence of these disorders among adults in their working years, likely due to life stressors and social factors often faced during this period (Lotzin et al., 2021).

To complement this general overview, Table 1 presents a detailed demographic breakdown by diagnosis, including gender distribution and descriptive age statistics. In both diagnostic categories, female patients represent the majority: 60.9% in Adjustment Disorder and 61.0% in Anxiety Disorder. Male representation is slightly higher in the Anxiety group (36.6%) compared to Adjustment Disorder (32.9%), while the proportion of patients with unknown gender is relatively low in both groups. Regarding age, the average for patients diagnosed with Adjustment Disorder is 44.4 years (SD = 18.3), and for Anxiety Disorder it is 42.7 years, with a median of 42 years in both cases. These findings confirm that the dataset is predominantly composed of middle-aged individuals, consistent across diagnostic categories, and reinforce the relevance of tailoring classification approaches to this demographic profile.

The age distribution provides valuable demographic insights and helps to contextualize the clinical data being analyzed, especially in terms of tailoring interventions for specific age groups. The relatively lower number of patients in the younger and older age ranges also raises important questions about the underrepresentation of these populations, possibly indicating a need for further exploration of psychiatric care in these demographics.

2.2 Machine Learning and Deep Learning models implemented

An intriguing research avenue was explored, focusing on the evaluation of ML and DL models to identify the most accurate approach for addressing the problem. Both linear and non-linear approaches were selected to determine which best suited textual data, given its high dimensionality and potential semantic

noise. Below, we describe the theoretical foundation and mathematical formulation of each model, along with the motivation for its selection.

- Random Forest: A versatile and widely used machine learning model that operates by constructing multiple decision trees during training and outputting the mode of the classes or the mean prediction of the individual trees. One key virtue of Random Forest in the context of clinical note classification, is its ability to handle high-dimensional and noisy data effectively, which is common in clinical settings (Al-Showarah et al., 2023). This robustness ensures reliable classification even when dealing with complex medical information as could be the Psychiatry EHRs, improving the model's accuracy and generalization on diverse clinical notes (Góngora Alonso et al., 2022). Mathematically, Random Forest combines multiple decision trees $h_i(x)$, where $h_i(x)$ represents the output of each individual tree i , and N is the total number of trees in the forest and the final prediction is obtained through majority voting for classification:

$$\hat{y} = \text{mode}\{h_1(x), h_2(x), \dots, h_N(x)\} \quad (1)$$

- Support Vector Classifier (SVC): a supervised machine learning model based on the Support Vector Machine (SVM) algorithm. It works by finding the optimal hyperplane that best separates the data into different classes. In this task, SVC tries to maximize the margin between the data points of different classes, which helps in achieving better generalization. SVC can capture complex relationships in the clinical notes, such as the nuanced patterns in clinical language (Elshehewy et al., 2023; Lyu et al., 2023). It is also robust to overfitting. The mathematical equation of the decision boundary is:

$$f(x) = w^T x + b \quad (2)$$

where w is the weight vector, x represents the input, and b is the bias term. The optimization process maximizes the margin $\frac{2}{\|w\|}$ subject to the following constraint where y_i represents the class label:

$$y_i(w^T x_i + b) \geq 1, \quad \forall i \quad (3)$$

- Decision Trees: A type of supervised learning algorithm that makes classifications based on a series of decision rules derived from the input features. The model works by recursively splitting the data into subsets based on feature values, creating branches that represent decision points. Each branch ultimately leads to a leaf node, which represents the predicted class or outcome. This model can help in classifying clinical notes because they are interpretable and can handle both numerical and categorical data. This interpretability is valuable in clinical settings, where understanding the reasoning behind a classification is important for trust and compliance (Vallée et al., 2023). Mathematically, node splitting is based on information gain or Gini index, defined below, where p_i is the proportion of instances of class i in dataset D .

$$Gini(D) = 1 - \sum_{i=1}^C p_i^2 \quad (4)$$

- XGBoost (eXtreme Gradient Boosting): is a powerful machine learning algorithm that is based on gradient boosting techniques. It works by creating an ensemble of decision trees, where each new tree corrects the errors made by the previous trees. The trees are added sequentially, with each one being optimized to reduce the total error. XGBoost uses gradient descent to minimize a loss function, which allows it to handle complex data patterns effectively (Mir and Sunanda, 2023). It can handle large datasets with high-dimensional features, such as the variety of terms and medical concepts found in clinical text. It also supports regularization, which helps prevent

330 overfitting—a common issue when working with detailed clinical data. Additionally, XGBoost
331 can efficiently process missing data and is relatively fast, making it well-suited for real-time or
332 large-scale applications in clinical settings (Ulhaq et al., 2023). Each tree in XGBoost is built by
333 minimizing the following regularized loss function, where $l(y_i, \hat{y}_i)$ represents the error function,
334 and $\Omega(f_k)$ penalizes model complexity to prevent overfitting:

$$L(\theta) = \sum_{i=1}^n l(y_i, \hat{y}_i) + \sum_k \Omega(f_k) \quad (5)$$

- 335 • SciBERT (Beltagy et al., 2019): An adapted version of BERT (Bidirectional Encoder Representations
336 from Transformers) that is specifically trained on scientific literature, including biomedical
337 and computer science articles, which makes it well-suited for handling the specialized language
338 in medical contexts. Its transformer-based architecture allows it to understand words in their full
339 context, making it effective for processing clinical notes. When fine-tuned for diagnostic classifica-
340 tion, SciBERT can accurately identify patterns in clinical text, recognizing terminology related to
341 various medical conditions (Tang et al., 2023). This makes it particularly valuable for automatically
342 categorizing clinical notes into diagnostic labels, improving the efficiency and accuracy of diagnosis
343 classification tasks in healthcare settings. The transformation function of each layer in BERT-based
344 models is given by the following formula where Q, K, V are the query, key, and value matrices,
345 respectively, and d_k is the key dimension:

$$z_i = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V \quad (6)$$

- 346 • DistilBERT (Sanh et al., 2019): is a distilled, or compressed, version of BERT that retains much
347 of BERT's effectiveness while being smaller, faster, and more efficient. It achieves this through a
348 process called 'knowledge distillation', where a smaller model (DistilBERT) is trained to mimic
349 the behavior of a larger model (BERT). DistilBERT has about 40% fewer parameters and is around
350 60% faster than BERT, but it retains around 97% of BERT's language understanding capabilities.
351 DistilBERT is useful for classifying clinical notes because it provides a good balance between
352 performance and computational efficiency (Abdelhalim et al., 2023). In clinical environments,
353 where there may be constraints on processing power or the need for quick responses, DistilBERT
354 can handle complex language and terminology effectively without requiring the resources that
355 full-sized BERT models do. This makes it suitable large-scale processing of clinical text, where
356 quick and accurate classifications are necessary (Oh et al., 2023; Le et al., 2023).

357 The selection of models in this study was driven by the need to evaluate both traditional machine
358 learning and deep learning approaches for classifying psychiatric clinical notes. Random Forest and
359 XGBoost were chosen for their strong generalization capabilities, while SVC was included to assess the
360 effectiveness of a linear decision boundary. Decision Tree was selected for its interpretability, which is
361 critical in clinical decision-making. In deep learning, SciBERT was used due to its training on biomedical
362 texts, making it well-suited for clinical language, while DistilBERT was included as a computationally
363 efficient alternative. This diverse set of models ensures a comprehensive evaluation of classification
364 techniques and performance.

365 **2.3 Hardware and software Specifications**

366 For the execution of all experiments, Jupyter Notebooks were used. These notebooks were run using
367 Python 3.9 and executed with the following hardware specifications: Intel(R) Core(TM) i7-9700K
368 CPU @ 3.60GHz, 32.0GB RAM, and an NVIDIA GeForce RTX 2080 graphics card. The code used
369 to perform the experiments described in this study is publicly available at the following repository:
370 <https://doi.org/10.5281/zenodo.14872650>.

3 EXPERIMENTS AND RESULTS

3.1 Data Preprocessing

Preprocessing clinical notes written in Spanish is a crucial step in preparing data for classification tasks using NLP techniques. Since clinical notes contain unstructured medical information, multiple cleaning and transformation techniques must be applied to enhance data quality and optimize the performance of machine learning models. The following preprocessing steps were performed in detail:

1. **Detection and handling of outliers:** During exploratory analysis, it was detected that clinical notes with fewer than 600 characters rarely contained relevant or substantial information. In many cases, they were administrative records without significant clinical data, and most of them lacked an actual diagnosis. To avoid including non-representative records, a minimum threshold of 600 characters was established in consultation with psychiatrists. By applying this criterion, only notes with sufficient clinical content were processed, ensuring more effective classification.
2. **Handling missing values:** Missing values in patient age and gender were addressed by extracting information from context using regular expressions. This process was carefully designed to improve the statistical reliability of the dataset and ensure a more accurate representation of the patient population.
3. **Lowercasing:** All text was converted to lowercase to reduce unnecessary variability and prevent the model from interpreting words with different casing as distinct entities. In medical language, some words may be capitalized due to writing conventions, but they are semantically equivalent to their lowercase counterparts. This normalization ensured a more homogeneous analysis and reduced the number of unique tokens in the model's vocabulary Chai (2022).
4. **Removal of special characters:** Accents were removed to ensure that a single word is not represented in multiple ways in the model and preserving only spaces and Spanish language characters. RAJESH and HITWARKAR (2023).
5. **Stopword removal:** Frequent words in Spanish that do not contribute meaningful information for classification, such as "el", "de", "que", "en", "un" and "una" were removed using a Spanish stop-word list adapted to the clinical context. Words like "patient," "symptom," or "treatment" were retained, as they are crucial in medical text analysis Sarica and Luo (2021).
6. **Lemmatization:** Lemmatization was applied using the spaCy library to reduce words to their canonical form, decreasing vocabulary dimensionality without losing meaning and maintaining the semantic integrity of clinical notes, which is crucial for understanding the context in medical text Babanejad et al. (2024).
7. **Removal of extra whitespaces:** Removing redundant whitespaces ensures text cleanliness and prevents models from misinterpreting the data as separate entities, improving tokenization accuracy Chai (2022).

This preprocessing pipeline optimized the representation of clinical notes, reducing data noise and improving the model's ability to capture key semantic patterns in medical texts.

3.2 Experimental Design

To evaluate the performance of various classification models in distinguishing between diagnoses of Adjustment Disorder and Anxiety Disorder, three different approaches were adopted for handling the training data. These approaches were: (1) without applying any oversampling techniques, (2) using Random Oversampling, and (3) employing the Synthetic Minority Over-sampling Technique (SMOTE).

Oversampling techniques, such as Random Oversampling and SMOTE, are commonly used in machine learning when dealing with imbalanced datasets, where one class is significantly underrepresented compared to the other. In this study, these techniques were explored to see how they could improve the model's ability to correctly classify both disorders, particularly the less common diagnosis, without overfitting to the majority class.

Additionally, to ensure that the distribution of classes (Adjustment Disorder and Anxiety Disorder) remained consistent in both the training and test sets, stratification was applied. Stratification is a method

420 that ensures the class proportions are maintained when splitting the dataset, which is particularly important
 421 in imbalanced datasets like this one. Without stratification, there is a risk that one of the sets (training or
 422 test) could have a disproportionate number of cases from one class, leading to unreliable performance
 423 metrics. By using stratified sampling, both the training (70%) and testing (30%) sets maintain the same
 424 distribution of Adjustment Disorder and Anxiety Disorder cases, providing a fair and consistent evaluation
 425 during model training and testing.

426 This step was essential for obtaining reliable performance measurements, as class imbalance can
 427 otherwise skew model performance toward the majority class, resulting in misleadingly high accuracy that
 428 does not reflect true generalization. Stratification helps prevent this by ensuring that both the minority and
 429 majority classes are well-represented in each dataset split, allowing the model to learn from a balanced
 430 representation of both diagnoses.

431 The classification models selected for this task were chosen for their varied approaches and capabilities
 432 in handling different types of data. These models included traditional machine learning models such as
 433 Random Forest, SVM, and Decision Tree, as well as more advanced models like XGBoost. In addition,
 434 two pre-trained transformer-based models, DistilBERT and SciBERT, were employed to leverage their
 435 capacity for understanding complex text patterns, particularly in the context of clinical notes.

436 Each model was evaluated based on two primary metrics: Accuracy and F1-Score. Accuracy provides
 437 a general measure of how often the model makes correct predictions and F1-Score gives a more balanced
 438 view of model performance in this context.

439 **3.3 Results**

440 This subsection describes the results of the experiments conducted on all models, both with and without
 441 the use of oversampling techniques. Table 2 presents the performance metrics for each model, highlighting
 442 their classification capabilities. The evaluation focuses on key metrics, particularly accuracy and F1-Score,
 443 to assess the effectiveness of the models under these conditions.

444 **3.3.1 Models without Oversampling Techniques**

445 The classification models were first evaluated without applying any oversampling techniques. The models
 446 demonstrated good performance, though there was significant variability among them.

447 The XGBoost model achieved the best results, with an accuracy of 96% and an F1-Score of 0.97,
 448 indicating excellent classification ability. The Decision Tree model followed, with an accuracy of 93%
 449 and an F1-Score of 0.94. These results suggest that tree-based models, particularly XGBoost, are highly
 450 effective for the task of classifying clinical notes in this dataset. The Random Forest model also showed
 451 satisfactory performance with an accuracy of 81% and an F1-Score of 0.87. However, the SVC model
 452 performed worse, with an accuracy of 70% and an F1-Score of 0.81, indicating that it struggled to
 453 effectively capture the relationships between features and classes in the data. The pre-trained transformer
 454 models (DistilBERT and SciBERT) performed similarly, both achieving an accuracy of 91% and an
 455 F1-Score of 0.93. This suggests that these language models, specialized in scientific and clinical text,
 456 are particularly useful for this task, outperforming simpler models like SVC and Random Forest. The
 457 results obtained without the application of oversampling techniques highlight the strong performance of
 458 the XGBoost and Decision Tree models, as well as the effectiveness of pre-trained transformer models.
 459 However, the SVC model showed limitations in its classification capability in this context.

| Exp | Metric | Rand.Forest | SVC | Dec.Tree | XGB | DistilBERT | SciBERT |
|-------|----------|-------------|------|----------|-------------|------------|---------|
| WO | Accuracy | 0.81 | 0.70 | 0.93 | 0.96 | 0.91 | 0.91 |
| | F1-Score | 0.87 | 0.81 | 0.94 | 0.97 | 0.93 | 0.93 |
| RO | Accuracy | 0.81 | 0.70 | 0.87 | 0.96 | 0.55 | 0.91 |
| | F1-Score | 0.87 | 0.81 | 0.90 | 0.97 | 0.55 | 0.93 |
| SMOTE | Accuracy | 0.87 | 0.70 | 0.88 | 0.96 | 0.91 | 0.91 |
| | F1-Score | 0.90 | 0.81 | 0.90 | 0.97 | 0.93 | 0.93 |

Table 2. Models Performance Across Experiments (Exp = Experiment, XGB = XGBoost, WO = Without Oversampling, RO = Random Oversampling, SMOTE)

460 3.3.2 Models with Random Oversampling

461 This subsection presents the performance of the models after applying random oversampling to balance
462 the dataset. The introduction of random oversampling had mixed effects on model performance.

463 The XGBoost model continued to achieve the highest performance, maintaining an accuracy of 96%
464 and an F1-Score of 0.97, consistent with the results without oversampling. This suggests that the XGBoost
465 model is robust to class imbalance, and the oversampling did not significantly alter its ability to classify
466 the clinical notes. The Decision Tree model saw a slight decrease in performance compared to the results
467 without oversampling. Its accuracy dropped from 93% to 87%, and the F1-Score decreased to 0.90. This
468 may suggest that random oversampling introduced some noise, reducing the model's ability to generalize
469 well to the test data. The Random Forest model showed no change in performance, with accuracy and
470 F1-Score remaining at 81% and 0.87, respectively. Similarly, the SVC model's performance remained
471 largely unchanged, with an accuracy of 70% and an F1-Score of 0.81. These results indicate that random
472 oversampling did not provide a substantial improvement for these models in this classification task.

473 Notably, the DistilBERT model experienced a significant drop in performance when random over-
474 sampling was applied. Its accuracy fell to 55%, and its F1-Score dropped to 0.55, suggesting that this
475 transformer-based model was negatively affected by the oversampling technique. On the other hand,
476 SciBERT maintained its strong performance, with an accuracy of 91% and an F1-Score of 0.93, indicating
477 that it was more resilient to the oversampling method. Random oversampling had varying effects on
478 model performance. While it did not lead to improvements in most models, XGBoost maintained its
479 high level of accuracy, and SciBERT remained effective. However, the significant drop in performance
480 for DistilBERT suggests that careful consideration is needed when applying oversampling techniques,
481 especially with transformer-based models.

482 3.3.3 Models with SMOTE

483 This subsection outlines the performance of the models after applying SMOTE to address class imbalance.
484 Compared to random oversampling, SMOTE generally had a more positive impact on model performance.
485 Once again, the XGBoost model achieved the highest accuracy of 96% and an F1-Score of 0.97, demon-
486 strating consistency across different data balancing techniques. This reinforces XGBoost's robustness and
487 adaptability to imbalanced datasets, as SMOTE did not significantly alter its performance. The Decision
488 Tree model showed a slight improvement with SMOTE compared to random oversampling, reaching
489 an accuracy of 88% and an F1-Score of 0.90. This marginal increase indicates that SMOTE helped the
490 model better generalize, although the performance is still lower than without any oversampling technique.
491 The Random Forest model also saw an improvement, with accuracy rising from 81% to 87% and the
492 F1-Score improving to 0.90. This suggests that SMOTE was more effective than random oversampling in
493 improving the model's ability to classify the minority class without overfitting to the majority class. SVC,
494 however, did not show any noticeable improvement, with its accuracy remaining at 70% and an F1-Score
495 of 0.81, similar to its performance without any oversampling technique. This indicates that SVC's ability
496 to capture relationships in the dataset was not enhanced by SMOTE.

497 For the transformer-based models, both DistilBERT and SciBERT maintained strong and consistent
498 performance, each achieving an accuracy of 91% and an F1-Score of 0.93. Unlike with random oversam-
499 pling, DistilBERT's performance remained stable with SMOTE, indicating that the synthetic examples
500 generated by this method may have been better aligned with the underlying data distribution, thereby
501 avoiding the performance degradation observed earlier.

502 SMOTE had a generally positive impact on model performance, especially for Random Forest and
503 Decision Tree, improving their ability to handle imbalanced data. XGBoost maintained its exceptional
504 performance, and the transformer models continued to show resilience, with DistilBERT recovering from
505 its previous drop in performance with random oversampling.

506 3.4 Hyperparameter Tuning

507 This subsection presents the results of hyperparameter tuning performed on all models, with and without
508 oversampling techniques, to optimize their performance. The complete hyperparameter search space
509 for each model is summarized in Table 3. This information provides a clearer view of the experimental
510 setup and supports the reproducibility of the results. A 3-fold cross-validation was applied during
511 hyperparameter search to ensure robust evaluation of each configuration. The results for each model after
512 tuning are shown in Table 4. The goal of hyperparameter tuning was to improve the classification metrics,
513 primarily focusing on accuracy and F1-Score.

| Model | Hyperparameters |
|---------------|--|
| Random Forest | n_estimators: [30, 97, 165, 232, 300]; max_features: ['sqrt', 'log2']; max_depth: [10, 20, 30, 40, 50, None]; min_samples_split: [2, 5, 10]; min_samples_leaf: [1, 2, 4]; bootstrap: [True, False] |
| SVM | C: [0.01, 0.1, 1, 2, 3, 4, 5, 10, 15, 50, 100, 1000]; gamma: [1, 0.1, 0.01, 0.001]; kernel: ['rbf', 'linear', 'sigmoid', 'poly'] |
| Decision Tree | criterion: ['gini', 'entropy', 'log_loss']; splitter: ['best', 'random']; max_depth: 1–29; min_samples_split: 1–19; min_samples_leaf: 1–19; max_features: ['sqrt', 'log2', None]; min_weight_fraction_leaf: [0.0]; random_state: [100] |
| XGBoost | objective: ['binary:logistic', 'binary:logitraw', 'binary:hinge']; learning_rate: [0.1, 0.3, 0.5]; n_estimators: [100, 200, 300, 400]; min_child_weight: [1, 5, 10]; gamma: [1, 2, 5]; subsample: [0.6, 0.8, 1.0]; colsample_bytree: [0.6, 0.8, 1.0]; max_depth: [2, 3, 4, 5] |
| DistilBERT | learning_rate: [1e-5, 3e-5, 5e-5]; batch_size: [8, 16, 32]; epochs: [3, 5, 10] |
| SciBERT | learning_rate: [1e-5, 3e-5, 5e-5]; batch_size: [8, 16]; epochs: [3, 5, 10] |

Table 3. Hyperparameter search space for each model.**3.4.1 Hyperparameter Tuning without Oversampling**

After tuning the hyperparameters, most models showed improved performance when no oversampling techniques were applied. Notably, the Decision Tree model experienced a significant boost, with accuracy rising from 93% to 96% and the F1-Score reaching 0.97. This suggests that fine-tuning the model parameters helped improve its capacity to better distinguish between the classes. The SVC model also demonstrated substantial improvements, with its accuracy increasing from 70% to 88% and its F1-Score reaching 0.91. These improvements reflect the positive impact of hyperparameter optimization on SVC's ability to better handle the complex relationships in the dataset. The Random Forest model improved slightly, with accuracy reaching 86% and an F1-Score of 0.90. Meanwhile, the XGBoost model saw a small decline in accuracy (from 96% to 93%) after hyperparameter tuning, though it still maintained exceptional performance. The slight decrease in performance might indicate that the default parameters were already close to optimal for this model. For the transformer models, both DistilBERT and SciBERT improved their accuracy to 96% and their F1-Scores to 0.97. These gains suggest that tuning transformer-specific parameters, such as learning rate and number of epochs, helped these models better capture the nuances in the clinical text, further boosting their effectiveness.

3.4.2 Hyperparameter Tuning with Random Oversampling

In the models trained with random oversampling, hyperparameter tuning led to notable improvements for the SVC model, which saw its accuracy rise to 88% and its F1-Score improve to 0.91, making it much more competitive compared to its previous performance. The Decision Tree model also benefited from

| Exp | Metric | Rand.Forest | SVC | Dec.Tree | XGB | DistilBERT | SciBERT |
|-------|----------|-------------|------|-------------|-------------|-------------|-------------|
| WO | Accuracy | 0.86 | 0.88 | 0.96 | 0.93 | 0.96 | 0.96 |
| | F1-Score | 0.90 | 0.91 | 0.97 | 0.94 | 0.97 | 0.97 |
| RO | Accuracy | 0.84 | 0.88 | 0.96 | 0.96 | 0.94 | 0.94 |
| | F1-Score | 0.89 | 0.91 | 0.97 | 0.97 | 0.95 | 0.95 |
| SMOTE | Accuracy | 0.83 | 0.88 | 0.91 | 0.96 | 0.96 | 0.96 |
| | F1-Score | 0.87 | 0.91 | 0.94 | 0.97 | 0.97 | 0.97 |

Table 4. Models Performance Across Experiments using Hyperparameter Tuning (Exp = Experiment, XGB = XGBoost, WO = Without Oversampling, RO = Random Oversampling, SMOTE)

533 tuning, achieving a significant boost in accuracy (96%) and F1-Score (0.97), indicating that the optimized
 534 parameters helped counterbalance the challenges posed by the oversampled data. The Random Forest
 535 model experienced a slight increase in performance after tuning, with accuracy reaching to 84% and the
 536 F1-Score to 0.89. XGBoost maintained its top performance, with both accuracy and F1-Score remaining
 537 at 96% and 97% respectively, further emphasizing its robustness to both data imbalance and parameter
 538 adjustments.

539 The transformer models, DistilBERT and SciBERT, both showed improvements with accuracy and
 540 F1-Scores rising to 94% and 0.95, respectively, indicating that the combination of random oversampling
 541 and tuning positively impacted their ability to classify the clinical notes accurately.

542 **3.4.3 Hyperparameter Tuning with SMOTE**

543 When SMOTE was used in conjunction with hyperparameter tuning, the results were similarly positive.
 544 The SVC model achieved an accuracy of 88% and an F1-Score of 0.91, consistent with its performance
 545 under other oversampling techniques. The Decision Tree model experienced a performance increase, with
 546 accuracy rising to 91% and an F1-Score of 0.94. Random Forest, however, showed a slight decrease in
 547 performance after tuning, with accuracy dropping to 83% and an F1-Score of 0.87, suggesting that tuning
 548 in combination with synthetic data did not favor this model. XGBoost continued to achieve excellent
 549 results, maintaining an accuracy of 96% and an F1-Score of 0.97. The transformer models, DistilBERT
 550 and SciBERT, also improved after tuning, with both achieving an accuracy of 96% and a F1-Score of
 551 0.97. Hyperparameter tuning was generally effective in enhancing model performance across various
 552 techniques. The Decision Tree and SVC models saw the most significant improvements, while XGBoost
 553 remained highly consistent. Transformer models also benefited notably from the optimization process.

554 **3.5 Computational Performance**

555 The computational time required for hyperparameter tuning varied significantly across the models, as
 556 summarized in Table 5. Traditional machine learning models such as Random Forest, SVM, Decision
 557 Tree, and XGB exhibited relatively low computational costs, with average times of 0.912 seconds
 558 (Random Forest), 0.091 seconds (SVM), 0.003 seconds (Decision Tree), and 0.103 seconds (XGB) per
 559 configuration.

560 In contrast, transformer-based models such as DistilBERT and SciBERT required substantially
 561 higher computational resources, with average tuning times of 75.70 seconds and 65.52 seconds per
 562 configuration, respectively. These results highlight a clear computational trade-off: while traditional
 563 models are significantly more efficient in hyperparameter tuning, transformer-based models demand
 564 considerably more processing time. However, prior studies suggest that this increased computational cost
 565 often translates into superior performance in terms of accuracy and generalization Benítez-Andrade et al.
 566 (2022); Meléndez et al. (2024).

567 **4 CONCLUSION AND DISCUSSION**

568 This research has contributed to the field of clinical text classification by examining the effectiveness
 569 of different machine learning models in distinguishing between patients diagnosed with Adjustment
 570 Disorder and Anxiety Disorder based on clinical notes. Several important findings emerged from this
 571 study, highlighting the strengths and limitations of the models employed, as well as the impact of applying
 572 oversampling techniques to address class imbalance in the dataset.

| Model | Combinations | Time (s) | Time/Combination |
|---------------|--------------|----------|------------------|
| Random Forest | 1080 | 985.267 | 0.912 |
| SVM | 192 | 17.482 | 0.091 |
| Decision Tree | 188442 | 651.747 | 0.003 |
| XGB | 11664 | 1204.597 | 0.103 |
| DistilBERT | 27 | 2043.965 | 75.70 |
| SciBERT | 18 | 1179.506 | 65.52 |

Table 5. Computational performance of each model used.

573 4.1 Model Performance

574 Among the models tested, XGBoost emerged as the best-performing algorithm, consistently demonstrating
 575 high accuracy and F1-Score across all experimental setups. Specifically, XGBoost achieved an F1-
 576 Score of 0.97 with and without the use of oversampling techniques, proving its robustness in handling
 577 the complexities of clinical text classification. The model maintained strong performance even after
 578 hyperparameter tuning, confirming its ability to effectively capture class distinctions while maintaining
 579 generalization, despite class imbalance in the dataset.

580 In contrast, the Support Vector Classifier (SVC) model exhibited the weakest performance, particularly
 581 without oversampling, where it struggled with an accuracy of 70% and an F1-Score of 0.81. This is likely
 582 due to the sensitivity of SVC to imbalanced datasets, where the minority class may be overshadowed
 583 by the majority class. Although hyperparameter tuning and oversampling techniques such as Random
 584 Oversampling and SMOTE improved SVC's performance (raising the F1-Score to 0.91 in some cases), its
 585 results remained below those of more advanced models like XGBoost, SciBERT, and DistilBERT. These
 586 findings indicate that while SVC can be a reliable option in certain domains, it may not be well-suited for
 587 imbalanced clinical text classification tasks without significant adjustments. Figures 4 and 5 illustrate the
 588 comparative performance of the machine learning and deep learning models, respectively.

589 4.2 The Impact of Oversampling Techniques

590 A key aspect of this research was the evaluation of two oversampling techniques: Random Oversampling
 591 and SMOTE. The results indicate that oversampling had a varying impact on model performance,
 592 particularly for models sensitive to class imbalance.

593 For models like Random Forest and XGBoost, Random Oversampling did not result in significant
 594 performance gains, and in some cases, even led to a slight drop in performance. For instance, DistilBERT
 595 experienced a considerable decline when random oversampling was applied, with the F1-Score dropping
 596 to 0.55. This suggests that Random Oversampling may introduce noise, particularly in more complex
 597 models, and thus does not consistently benefit all models.

598 SMOTE, on the other hand, proved to be a more effective technique for improving performance
 599 across various models. In particular, SMOTE enhanced the performance of models like Decision Tree and
 600 Random Forest, which achieved F1-Scores of 0.90 and 0.90, respectively, when applied. Furthermore,
 601 models such as XGBoost and transformer-based models like DistilBERT and SciBERT maintained their
 602 strong performance with SMOTE, both achieving F1-Scores of 0.97. The results indicate that SMOTE
 603 helped these models create more balanced decision boundaries without duplicating existing data points,
 604 leading to more robust classification outcomes.

605 It was found that, while oversampling techniques generally improved performance, SMOTE was more
 606 effective across a range of models, particularly for complex architectures like XGBoost and transformer-
 607 based models.

608 4.3 Comparison Between Transformer Models and Traditional Machine Learning

609 The transformer-based models, DistilBERT and SciBERT, demonstrated strong results throughout the
 610 experiments, confirming their potential for natural language processing tasks in the healthcare domain. In
 611 comparison to traditional machine learning models such as Random Forest and SVC, the transformers-
 612 based models were better at capturing the nuances of clinical language, particularly when no oversampling
 613 techniques were applied.

614 SciBERT, pretrained on scientific texts, was particularly noteworthy, achieving an F1-Score of 0.97
 615 with SMOTE, highlighting its strength in parsing and classifying the specialized terminology found in




Figure 4. Visual comparison of accuracy and F1-score across machine learning models using different configurations (WO = Without Oversampling, RO = Random Oversampling, H = Hyperparameters)

616 clinical notes. DistilBERT performed well across all setups, reinforcing the potential of transformer-based
 617 models in healthcare text classification tasks. Notably, these transformer models remained highly effective,
 618 even when class imbalance was not addressed through oversampling techniques.

619 **4.4 Impact of Hyperparameter Tuning**

620 Hyperparameter tuning was a critical component in this study, as it helped optimize the performance of all
 621 the machine learning models. The results clearly show that hyperparameter tuning had a significant impact
 622 on improving classification metrics, particularly for models that initially struggled with imbalanced data
 623 or suboptimal settings.

624 The most notable improvement was observed in the Decision Tree model, where hyperparameter
 625 tuning increased its accuracy from 93% to 96% and its F1-Score to 0.97 when no oversampling was
 626 applied. This demonstrates that tuning allowed the Decision Tree model to make better splits and
 627 generalize more effectively on the data, leading to performance that matched the top-performing models
 628 such as XGBoost.

629 Similarly, the SVC model benefited substantially from hyperparameter tuning. Initially, SVC struggled
 630 with imbalanced data, but after tuning, its accuracy increased to 88% and its F1-Score improved to 0.91.
 631 These improvements indicate that carefully optimizing parameters like the kernel and gamma allowed
 632 SVC to better distinguish between the diagnostic categories.

633 The transformer models, DistilBERT and SciBERT, also saw improvements with hyperparameter
 634 tuning. Both DistilBERT and SciBERT achieved an accuracy of 96% and an F1-Score of 0.97 after
 635 tuning. These results suggest that while the transformers performed well without significant tuning, the




Figure 5. Visual comparison of accuracy and F1-score across deep learning models using different configurations (WO = Without Oversampling, RO = Random Oversampling, H = Hyperparameters)

636 fine-tuning of parameters like learning rate and number of epochs still provided marginal performance
637 boosts.

638 For the XGBoost model, however, hyperparameter tuning led to a slight reduction in accuracy,
639 dropping from 96% to 93%, although its F1-Score remained high at 0.94. This suggests that XGBoost
640 may have already been operating near its optimal settings.

641 Hyperparameter tuning proved to be a valuable step in improving model performance. While simpler
642 models like Decision Tree and SVC saw the most pronounced benefits, even advanced models such as
643 transformers and XGBoost showed gains in certain metrics, reaffirming the importance of hyperparameter
644 optimization in machine and deep learning workflows.

645 **4.5 Limitations and Future Directions**

646 Despite the strong performance of the models tested, several limitations should be acknowledged. First,
647 the dataset, although preprocessed, might still contain noise inherent to clinical notes, such as inconsistent
648 terminology or incomplete information, which could affect model performance. Future studies could
649 focus on refining the preprocessing pipeline to handle these nuances more effectively, potentially leading
650 to further improvements in classification accuracy.

651 Additionally, while this study demonstrated the value of oversampling techniques, there are alternative
652 methods for addressing class imbalance that were not explored, such as cost-sensitive learning or under-
653 sampling methods, which could be examined in future research. These techniques might offer more
654 efficient solutions, especially in scenarios where oversampling introduces overfitting or data redundancy.
655 Future directions could also consider exploring Variational Autoencoders (VAEs) as a generative approach
656 for oversampling.

657 Another avenue for future research involves evaluating additional classification models beyond those
658 tested in this study. Exploring more advanced deep learning architectures or novel transformer-based
659 models could further enhance classification performance, particularly in complex diagnostic scenarios.
660 Moreover, expanding the dataset to include clinical notes from patients presenting similar symptomatology
661 but ultimately receiving different diagnoses would provide a more challenging and realistic classification
662 setting. This would help assess the models' ability to capture subtle clinical distinctions, which is critical
663 in psychiatric evaluation.

664 Finally, although transformer models performed well, their computational cost and the need for large
665 datasets for fine-tuning present practical challenges. Future work could explore the use of more efficient
666 transformer architectures or hybrid models that combine the strengths of transformers and traditional
667 machine learning approaches.

668 REFERENCES

- 669 Abdelhalim, N., Abdelhalim, I., and Batista-Navarro, R. (2023). Training models on oversampled data
670 and a novel multi-class annotation scheme for dementia detection. In Naumann, T., Ben Abacha, A.,
671 Bethard, S., Roberts, K., and Rumshisky, A., editors, *Proceedings of the 5th Clinical Natural Language
672 Processing Workshop*, pages 118–124, Toronto, Canada. Association for Computational Linguistics.
- 673 Adekkattu, P., Furmanchuk, A., Wu, Y., Pathak, A., Patra, B. G., Bost, S., Morrow, D., Wang, G. H.-M.,
674 Yang, Y., Forrest, N. J., Luo, Y., Walunas, T. L., Lo-Ciganic, W., Gelad, W., Bian, J., Bao, Y., Weiner,
675 M., Oslin, D., and Pathak, J. (2024). Deep learning for identifying personal and family history of
676 suicidal thoughts and behaviors from EHRs. *npj Digital Medicine*, 7(1):260.
- 677 Al-Showarah, S., Al-Taie, A., Salman, H. E., Alzyadat, W., and Alkhalaileh, M. (2023). Predicting quality
678 medical drug data towards meaningful data using machine learning. *International Journal of Advanced
679 Computer Science and Applications*, 14(8).
- 680 Alanazi, A. (2023). Clinicians' Views on Using Artificial Intelligence in Healthcare: Opportunities,
681 Challenges, and Beyond. *Cureus*, 15(9):e45255. Place: United States.
- 682 ALSAGRI, H. S. and YKHLEF, M. (2020). Machine learning-based approach for depression detection
683 in twitter using content and activity features. *IEICE Transactions on Information and Systems*,
684 E103.D(8):1825–1832.
- 685 Babanejad, N., Davoudi, H., Agrawal, A., An, A., and Papagelis, M. (2024). The role of preprocessing for
686 word representation learning in affective tasks. *IEEE Transactions on Affective Computing*, 15(1):254–
687 272.
- 688 Beltagy, I., Lo, K., and Cohan, A. (2019). SciBERT: A pretrained language model for scientific text. In
689 Inui, K., Jiang, J., Ng, V., and Wan, X., editors, *Proceedings of the 2019 Conference on Empirical
690 Methods in Natural Language Processing and the 9th International Joint Conference on Natural
691 Language Processing (EMNLP-IJCNLP)*, pages 3615–3620, Hong Kong, China. Association for
692 Computational Linguistics.
- 693 Benítez-Andrade, J. A., Alija-Pérez, J.-M., Vidal, M.-E., Pastor-Vargas, R., and García-Ordás, M. T.
694 (2022). Traditional machine learning models and bidirectional encoder representations from transformer
695 (bert)–based automatic classification of tweets about eating disorders: Algorithm development and
696 validation study. *JMIR Med Inform*, 10(2):e34492.
- 697 Briganti, G. and Le Moine, O. (2020). Artificial intelligence in medicine: Today and tomorrow. *Frontiers
698 in Medicine*, 7.
- 699 Casey, P. and Bailey, S. (2011). Adjustment disorders: the state of the art. *World Psychiatry*, 10(1):11–18.
- 700 Chai, C. P. (2022). Comparison of text preprocessing methods. *Natural Language Engineering*, 29:509 –
701 553.
- 702 Cortes-Briones, J. A., Tapia-Rivas, N. I., D'Souza, D. C., and Estevez, P. A. (2022). Going deep into
703 schizophrenia with artificial intelligence. *Schizophrenia Research*, 245:122–140. Computational
704 Approaches to Understanding Psychosis.
- 705 Dobbins, N. J., Chipkin, J., Byrne, T., Ghabra, O., Siar, J., Sauder, M., Huijon, R. M., and Black, T. M.
706 (2024). Deep learning models can predict violence and threats against healthcare providers using
707 clinical notes. *npj Mental Health Research*, 3(1):61.
- 708 Elshewey, A. M., Shams, M. Y., El-Rashidy, N., Elhadly, A. M., Shohieb, S. M., and Tarek, Z. (2023).
709 Bayesian optimization with support vector machine model for parkinson disease classification. *Sensors*,
710 23(4).
- 711 European Parliament (2016). Regulation (eu) 2016/679 of the european parliament and of the council of
712 27 april 2016 on the protection of natural persons with regard to the processing of personal data and on
713 the free movement of such data, and repealing directive 95/46/ec (general data protection regulation).
- 714 García-Barragán, A., González Calatayud, A., Solarte-Pabón, O., Provencio, M., Menasalvas, E., and
715 Robles, V. (2024). GPT for medical entity recognition in Spanish. *Multimedia Tools and Applications*.
- 716 Góngora Alonso, S., Marques, G., Agarwal, D., De la Torre Díez, I., and Franco-Martín, M. (2022).
717 Comparison of machine learning algorithms in the prediction of hospitalized patients with schizophrenia.
718 *Sensors*, 22(7).
- 719 Henriques, G. and Michalski, J. (2020). Defining behavior and its relationship to the science of psychology.
720 *Integrative Psychological and Behavioral Science*, 54(2):328–353.
- 721 Kaka, H., Michalopoulos, G., Subendran, S., Decker, K., Lambert, P., Pitz, M., Singh, H., and Chen,
722 H. (2022). Pretrained Neural Networks Accurately Identify Cancer Recurrence in Medical Record.

- 723 *Studies in health technology and informatics*, 294:93–97. Place: Netherlands.
- 724 Kendler, K. S., Zachar, P., and Craver, C. (2011). What kinds of things are psychiatric disorders?
725 *Psychological Medicine*, 41(6):1143–1150.
- 726 Kodipalli, A. and Devi, S. (2023). Analysis of fuzzy based intelligent health care application system for
727 the diagnosis of mental health in women with ovarian cancer using computational models. *Intelligent
728 Decision Technologies*, 17(1):31–42. Publisher: IOS Press.
- 729 Le, T.-D., Jouvet, P., and Noumeir, R. (2023). A small-scale switch transformer and nlp-based model for
730 clinical narratives classification. *ArXiv*, abs/2303.12892.
- 731 Li, Z., Hu, Y., Lane, S., Selek, S., Shahani, L., Machado-Vieira, R., Soares, J., Xu, H., Liu, H., and Huang,
732 M. (2024). Suicide phenotyping from clinical notes in safety-net psychiatric hospital using multi-label
733 classification with pre-trained language models.
- 734 Liu, R., Rong, Y., and Peng, Z. (2020). A review of medical artificial intelligence. *Global Health Journal*,
735 4(2):42–45.
- 736 Lotzin, A., Krause, L., Acquarini, E., Ajduković, D., Ardino, V., Arnberg, F., Böttche, M., Bragesjö, M.,
737 Dragan, M., Figueiredo-Braga, M., Gelezelyte, O., Grajewski, P., Anastassiou-Hadjicharalambous,
738 X., Javakhishvili, J., Kazlauskas, E., Lenferink, L., Lioupi, C., Lueger-Schuster, B., Tsiskarishvili, L.,
739 Mooren, T., Sales, L., Stevanović, A., Zrnić, I., Schäfer, I., and Consortium, A. S. (2021). Risk and
740 protective factors, stressors, and symptoms of adjustment disorder during the covid-19 pandemic – first
741 results of the estss covid-19 pan-european adjust study. *European Journal of Psychotraumatology*, 12.
- 742 Lyu, Y., Xu, Q., Yang, Z., and Liu, J. (2023). Prediction of patient choice tendency in medical decision-
743 making based on machine learning algorithm. *Frontiers in Public Health*, 11.
- 744 Meléndez, R., Ptaszynski, M., and Masui, F. (2024). Comparative investigation of traditional machine-
745 learning models and transformer models for phishing email detection. *Electronics*, 13(24).
- 746 Midasala, V. D., Prabhakar, B., Krishna Chaitanya, J., Sirnivas, K., Eshwar, D., and Kumar, P. M.
747 (2024). Mfeuslnet: Skin cancer detection and classification using integrated ai with multilevel feature
748 extraction-based unsupervised learning. *Engineering Science and Technology, an International Journal*,
749 51:101632.
- 750 Mir, S. and Sunanda (2023). “heart disease prediction and severity level classification”: A machine
751 learning approach with feature selection technique. *2023 14th International Conference on Computing
752 Communication and Networking Technologies (ICCCNT)*, pages 1–7.
- 753 Moitra, M., Santomauro, D., Collins, P. Y., Vos, T., Whiteford, H., Saxena, S., and Ferrari, A. J. (2022).
754 The global gap in treatment coverage for major depressive disorder in 84 countries from 2000–2019: A
755 systematic review and bayesian meta-regression analysis. *PLOS Medicine*, 19(2):1–16.
- 756 Msosa, Y. J., Grauslys, A., Zhou, Y., Wang, T., Buchan, I., Langan, P., Foster, S., Walker, M., Pearson, M.,
757 Folarin, A., Roberts, A., Maskell, S., Dobson, R., Kullu, C., and Kehoe, D. (2023). Trustworthy data
758 and ai environments for clinical prediction: Application to crisis-risk in people with depression. *IEEE
759 Journal of Biomedical and Health Informatics*, 27(11):5588–5598.
- 760 Nemesure, M. D., Heinz, M. V., Huang, R., and Jacobson, N. C. (2021). Predictive modeling of
761 depression and anxiety using electronic health records and a novel machine learning approach with
762 artificial intelligence. *Scientific Reports*, 11(1):1980.
- 763 Nimavat, N., Hasan, M. M., Mandala, G., Singh, S., Bhangu, R., and Bibi, S. (2023). Mortality Rate in
764 Schizophrenia. In Chatterjee, I., editor, *Cognizance of Schizophrenia:: A Profound Insight into the
765 Psyche*, pages 303–312. Springer Nature Singapore, Singapore.
- 766 Oh, J., Kim, M., Park, H., and Oh, H. (2023). Are you depressed? analyze user utterances to detect
767 depressive emotions using distilbert. *Applied Sciences*.
- 768 Park, J.-H., Shin, Y.-B., Jung, D., Hur, J.-W., Pack, S. P., Lee, H.-J., Lee, H., and Cho, C.-H. (2025).
769 Machine learning prediction of anxiety symptoms in social anxiety disorder: utilizing multimodal data
770 from virtual reality sessions. *Frontiers in Psychiatry*, 15.
- 771 Quanyang, W., Yao, H., Sicong, W., Linlin, Q., Zewei, Z., Donghui, H., Hongjia, L., and Shijun, Z. (2024).
772 Artificial intelligence in lung cancer screening: Detection, classification, prediction, and prognosis.
773 *Cancer Medicine*, 13(7):e7140. e7140 CAM4-2023-11-5724.R1.
- 774 RAJESH, M. A. and HIWARKAR, D. T. (2023). Exploring preprocessing techniques for natural lan-
775 guagetext: A comprehensive study using python code. *international journal of engineering technology
776 and management sciences*.
- 777 Rao, K. N., Arora, R. D., Dange, P., and Nagarkar, N. M. (2023). NLP AI Models for Optimizing Medical

- 778 Research: Demystifying the Concerns. *Indian Journal of Surgical Oncology*, 14(4):854–858.
- 779 Rubio-Martín, S., García-Ordás, M. T., Bayón-Gutiérrez, M., Prieto-Fernández, N., and Benítez-Andrade,
780 J. A. (2024). Enhancing ASD detection accuracy: a combined approach of machine learning and deep
781 learning models with natural language processing. *Health Information Science and Systems*, 12(1):20.
- 782 Sanh, V., Debut, L., Chaumond, J., and Wolf, T. (2019). Distilbert, a distilled version of bert: smaller,
783 faster, cheaper and lighter. *ArXiv*, abs/1910.01108.
- 784 Sarica, S. and Luo, J. (2021). Stopwords in technical language processing. *PLOS ONE*, 16(8):1–13.
- 785 Tang, X., Tran, A., Tan, J., and Gerstein, M. B. (2023). Gersteinlab at medica-chat 2023: Clinical note
786 summarization from doctor-patient conversations through fine-tuning and in-context learning. *arXiv
787 preprint*, 2305:546–554.
- 788 Ulhaq, S., Khan, G. Z., Ulhaq, I., Ullah, I., Rabbi, F., Zaman, G., and Khan (2023). Epilepsy seizures
789 classification with eeg signals: A machine learning approach. *Journal of Computer Science and
790 Technology Studies*.
- 791 Vallée, R., Vallée, J., Guillevin, C., Lallouette, A., Thomas, C., Rittano, G., Wager, M., Guillevin, R.,
792 and Vallée, A. (2023). Machine learning decision tree models for multiclass classification of common
793 malignant brain tumors using perfusion and spectroscopy mri data. *Frontiers in Oncology*, 13.
- 794 Wang, S., Sun, X., Li, X., Ouyang, R., Wu, F., Zhang, T., Li, J., and Wang, G. (2023). Gpt-ner: Named
795 entity recognition via large language models.
- 796 World Health Organization (2019). International statistical classification of diseases and related health
797 problems 10th revision (icd-10).
- 798 World Health Organization (2021). *Mental health atlas 2020*. World Health Organization.
- 799 World Health Organization (2022a). Mental disorders.
- 800 World Health Organization (2022b). Mental health and covid-19: Early evidence of the pandemic's
801 impact: Scientific brief.
- 802 Wu, K. (2024). Optimizing diabetes prediction with machine learning: Model comparisons and insights.
803 *Journal of Science & Technology*, 5(4):41–51.
- 804 Zhang, F., Geng, J., Zhang, D.-G., Gui, J., and Su, R. (2023). Prediction of cancer recurrence based on
805 compact graphs of whole slide images. *Computers in Biology and Medicine*, 167:107663.

Figure 1

Representation of the mean length of the clinical notes for each patient.

The y-axis indicates the mean length of the clinical notes, while the x-axis corresponds to patient IDs. The green-shaded region highlights the 95% of clinical notes that are above the applied threshold, whereas the red-shaded area represents the lower 5%




Figure 2

Representation of the prompt structure

The 'System' role defines the model's function, ensuring it adopts an expert perspective in psychiatry and entity recognition. Also it provides step-by-step guidance on how the model should process clinical notes, including handling special cases. The 'User' role supplies clinical notes as plain text, while the 'Assistant' role generates structured responses. A correct answer example is provided after the first clinical note to guide the model's output format, ensuring consistency and accuracy in entity extraction.




Figure 3

Demographic data from the patients found in the clinical notes

The pie chart shows the gender distribution of patients: women (orange), men (blue), and unknown (gray). The histogram represents the age distribution, with the X-axis indicating age and the Y-axis showing the number of clinical notes.




Figure 4

Machine Learning Models Performance Evolution

The x-axis represents different techniques: Without Oversampling (WO), Random Oversampling (RO), SMOTE, and their combinations with hyperparameter tuning (H). The y-axis shows the performance scores, with accuracy (yellow) and F1-score (red) used as evaluation metrics.




Figure 5

Deep Learning Models Performance Evolution

The x-axis represents different techniques: Without Oversampling (WO), Random Oversampling (RO), SMOTE, and their combinations with hyperparameter tuning (H). The y-axis shows the performance scores, with accuracy (yellow) and F1-score (red) used as evaluation metrics.

